

Start here.

One accessible view of the selected product direction, launch plan, naming decision, current market position, repository controls, and Matt's first-case-study opportunity radar.

Product direction Radar Setup to Brief to Review to native Pursuit

Brand candidate Worthward advances to validation, not launch

First case study Going ranks first at 85/100, pending proof

Current challenger Tebra ranks second at 84/100, pending hard application-gate proof

The product decision

Build an independent career decision system, not another job board, application-volume tool, or generic AI matcher. The recurring value is maintaining a user-approved Career Baseline, verifying opportunity integrity, mapping requirements to source-backed proof, comparing career economics and pursuit effort, and preparing external action only after approval.

Selected customer journey

1. Radar Setup defines what should be monitored and when the user should be interrupted.
2. Each dated Radar Brief explains the few opportunities that deserve attention or confirms that none do.
3. Opportunity Review separates proven, plausible, missing, risky, and excluded evidence.
4. The Decision Receipt records the read-only recommendation, main risk, user decision, and one next action.
5. A native Pursuit owns preparation; external-action approval remains separate and exact.

Defensible wedge

The platform wins by making the entire decision contract inspectable: baseline, job integrity, proof lineage, career economics, uncertainty, effort, user decision, external-action approval, and calibration. Scanning and alerts alone are not enough to justify a subscription.

The three full visual concepts appear at the end of this PDF. Detailed source files are linked from [the working hub](#).

One connected experience

SURFACE	DECISION	PURPOSE
Radar Setup	New screen required	Define alert criteria, material-upside threshold, cadence, notifications, and source health.
Radar Brief	Use Option 2	Make one dated result and honest no-action weeks easy to understand.
Opportunity Review	Use Option 3	Show every role requirement beside candidate evidence, state, source, and correction path.
Ranked row and Decision Receipt	Borrow from Option 1	Use components only; do not build the overloaded concept as a whole screen.
Pursuit Application	New screen required	Review the exact package and keep employer-site action separately approved and verifiable.

Theme and accessibility

Settings must offer System, Light, and Dark. System is the default. Both themes must preserve the same content, hierarchy, states, keyboard operation, visible focus, semantic structure, contrast, loading, empty, error, conflict, and no-opportunity behavior. A dark theme is not a simple inversion of the light theme.

Truth corrections before implementation

- Going fit is 85/100, not 86/100.
- Published compensation starts at \$175,000 plus equity; the public upper bound is unknown.
- The decision remains Investigate until proof validation.
- The job post proves an employer requirement, not Matt's candidate evidence.
- The user's decision appears once as Pursue, Watch, or Pass.
- Remove fixed public duration claims such as an eight-minute review.

[Open the full visual-direction decision.](#)

[Open the complete 21-screen system and interaction contract.](#)

[Open the theme, responsive, and accessibility specification.](#)

Implementation acceptance boundary

The concept images are design evidence, not implementation QA. Approval still requires responsive reflow, keyboard and screen-reader operation, accessible proof states, provenance, theme parity, realistic system states, and matched-viewport visual comparison.

Working packets

PACKET	WHAT IT CONTROLS
Product authority	Source hierarchy, Career Baseline, opportunity integrity, proof states, scores, approvals, state machine, UX rules, and release gates.
Platform screen system	All 21 screen families, routes, Radar object model, Pursuits, concept mapping, interaction semantics, responsive rules, and release sequencing.
Business launch and validation	Beachhead, concierge-assisted paid beta, offer ladder, economics tests, observed outcomes, decision gates, and stop conditions.
Discovery and paid-beta script	Neutral discovery, live decision test, competitor and willingness-to-pay contrast, no-opportunity-week test, beta invitation, and observation sheet.
Data, telemetry, and tests	Source-to-decision model, events, metrics, corrections, calibration, uncertainty, approvals, and acceptance tests.
Governance and reset	Delivery controls, source-of-truth reconciliation, ownership, WIP, and implementation boundaries.
Personas and journey	The first paid beachhead and the path from career baseline to recurring radar to pursuit.

Monetization sequence

1. Demonstrate one Career Baseline diagnostic or Decision Receipt.
2. Validate judgment through a concierge-assisted paid beta before engineering continuous monitoring.
3. Charge recurring fees for maintained, verified decisions, not alerts alone.
4. Offer a higher-value Pursuit Sprint when an opportunity clears the evidence and economics bar.
5. Use observed decisions, corrections, applications, interviews, offers, and no-action weeks to calibrate the product without inventing causal claims.

What the paid beta must prove

- Users change or defend a real career decision because of the Decision Receipt.
- Users value verified restraint during a no-opportunity week.
- The product reduces wasted pursuits without suppressing credible upside.
- Users will pay for independent judgment despite free and low-cost scanning alternatives.
- The recurring service remains useful after a user gets a new job.

The moat is the inspectable decision contract

Always-on scanning, job radar, ranked matches, career intelligence, reason explanations, resume generation, and one next action are active competitor claims. The product must make this combination observable:

- 1. user-approved Career Baseline,
- 2. canonical employer, job, freshness, and duplicate integrity,
- 3. versioned proof lineage for every externally usable claim,
- 4. career economics compared with the user's current position,
- 5. visible uncertainty and pursuit effort,
- 6. a portable Decision Receipt,
- 7. separate user decision and external-action approval,
- 8. calibration from corrections and observed outcomes.

[Open the current competition delta.](#)

Brand-name decision

RANK	NAME	STATUS	MAIN RISK
1	Worthward	Advance to validation	Heard as Worthword; ward associations; possible personal-worth ambiguity.
2	ConsiderNext	Functional backup	Generic, less ownable, and can sound like an unfinished instruction.
3	Nextworthy	Emotional-safety backup	Can imply that the person, rather than the opportunity, is being judged.

Worthward must pass hear-once spelling, five-second category comprehension, emotional safety, real-message comprehension, delayed recall, and formal legal clearance. If it fails, run another round rather than forcing a weaker backup.

[Open the full naming decision and validation protocol.](#)

Going stays first; Tebra is the close challenger

RANK	ROLE	SCORE	DECISION	OPEN PROOF
1	Going, Director, Lifecycle Marketing	85	Pursue conditionally	Braze depth, push/in-app, lifecycle analytics, and exact leadership scope.
2	Tebra, Director, GTM Technology	84	Pursue conditionally	High-velocity Sales/SDR accountability, pipeline cadence, Salesforce and GTM tools, and ops-team composition.
3	Pendulum, Senior Director, Lifecycle Marketing/Retention	83	Pursue selectively	Subscription-retention leadership, cohort/LTV analysis, SQL partnership, and push/in-app.

Going remains first because its risk is narrower and more inspectable. Tebra should not displace it solely because the title and published Zone 1 salary band are more attractive.

Evidence work ready for Matt

- [Going proof prefill](#): source-safe evidence plus seven only-Matt questions.
- [Tebra proof prefill](#): every live application gate, safe evidence, prohibited inferences, and seven only-Matt questions.
- [Fresh Opportunity Radar](#): eight ranked roles with freshness, integrity, economics, proof risk, and pursuit decision.

Current next evidence gate

Matt's seven Going answers determine whether the first case-study pursuit can advance. Short fragments, `none`, and `I do not remember` are valid. The system must become more accurate when evidence is missing, not more confident.

Job Filter repository decision

Fetch origin/main is the clean future base. The current branch is quarantined as evidence because it is 11 commits behind and 3 ahead, includes a dirty bootstrap file, and contains a mixed 165-file commit. Preserve and review additive material; never merge or cherry-pick the mixed commit wholesale.

[Open the cleanup and salvage plan.](#)

State-change receipt


Teal records mutated	No
ClickUp records mutated	No
Applications submitted	No
Outreach sent	No
Product implementation claimed	No
Repository cleanup executed	No. The evidence, preservation, and source-of-truth plan are complete; destructive or overlapping cleanup remains gated.

Capacity continuity

The earlier selected-model capacity notice was transient. Work continued from persisted files and the existing goal. The durable working hub, source packets, role audits, and repo evidence mean a future model interruption should not require restarting the research.

Last updated 2026-07-17. This working pack separates verified facts, inferences, missing evidence, product decisions, and external-action approval. It is not an efficacy claim, legal clearance, or application authorization.

#1

BEST MATCH

86 /100

Overall score

Director of Lifecycle Marketing

Pursue after proof check

Published Jul 10 · 6 days old

Remote · United States

\$175k–\$190k + equity

Canonical source verified

Best current match for your lifecycle systems, Braze, CRM scale, and remote compensation target.

Lane fit

90

Excellent

Mandate fit

85

Strong

Evidence

75

Good

Compensation

80

Meets target

Freshness

90

Very recent

Why it ranks first

- Lifecycle systems supporting 600K+ customers

Proof found

Job post cites 600K+ active members powered by lifecycle programs.
- Braze + CRM audience growth to 250K+

Proof found

Role requires Braze and growing CRM audience to 250K+.
- Remote base range clears your \$180K target

Proof found

\$175K–\$190K base + equity aligns with your compensation target.

Top alternatives (worthy, but score lower)

- G2

VP Revenue Marketing

Investigate

Score 72

Published Jul 8
- Upstart

Director Lifecycle

Monitor

Score 68

Published Jul 7

Decision receipt



Investigate 3 proof gaps

- Proven (4)

600K+ lifecycle scaleJob post

Braze experienceJob post

250K+ CRM audienceJob post

Team leadershipJob post
- Plausible (2)

End-to-end lifecycle ownershipJob post

Behavioral segmentationJob post
- Missing (3)

Braze Canvas / Liquid depthNot found

Push + in-app ownershipNot found

Cohort / LTV modelingNot found
- Excluded (1)

Unsupported churn or LTV metricsExcluded

Main risk

The role may expect deeper platform mechanics than your verified record currently proves.

Verify 3 proof gaps

HoldRejectCorrect evidence

Source

Going Careers

Verification

Canonical

Your opportunity brief

Thursday, July 16, 2026 • 14 roles reviewed • 3 deserve attention

WORTH YOUR ATTENTION

- 1

GOING

Going

Director of Lifecycle Marketing

\$175k–\$190k • 6 days old

86

Investigate
- 2

G2

G2

VP Revenue Marketing

\$190k–\$210k • 4 days old

72

Investigate
- 3

Up

Upstart

Director Lifecycle

\$165k–\$185k • 3 days old

68

Investigate

FILTERED OUT THIS WEEK

🚫

11 low-confidence or stale roles suppressed

See why

TOP OPPORTUNITY

Going could be your best current move, if three proof gaps hold up.

Going • Director of Lifecycle Marketing • Remote US

📅

Published Jul 10 • 6 days old

💰

\$175k–\$190k + equity

🛡️

Canonical source verified

🎯

Fit 86/100

Going operates large-scale lifecycle programs with Braze, a 250K+ CRM audience, and a systems-first approach to personalization that aligns with your background in lifecycle strategy and cross-functional leadership. The compensation range clears your \$180K target.

👤

Why now

📈

600K+ lifecycle scale

EVIDENCE

👥

250K+ CRM audience

EVIDENCE

💵

Remote range clears \$180K target

EVIDENCE

⚠️ What could break the fit

- 1

Canvas / Liquid depth

Unknown depth with email templates and landing page tooling.

PROVENANCE: JOB DESCRIPTION
- 2

Push + in-app ownership

No explicit ownership of push or in-app channels.

PROVENANCE: JOB DESCRIPTION
- 3

Cohort / LTV modeling

No mention of cohort analysis or LTV-driven decisioning.

PROVENANCE: JOB DESCRIPTION

Your decision

✓

Pursue

⏸

Hold

✗

Reject

🔍

Investigate

🚫

Ignore

Review the 3 proof gaps

>

Hold

Reject



Going

Director of Lifecycle Marketing

86 / 100

Investigate before applying

Remote • United States

\$175,000 – \$190,000 + equity

Published Jul 10 • 6 days old

Canonical source verified

Strong lifecycle-systems match. Three missing proof details determine whether this is a credible \$180K+ pursuit.

What the role requires vs. what you can prove

Requirement	Evidence State	Your Proof	Source
Lifecycle programs at scale	PROVEN	Systems supporting 600K+ customers	Prosper / verified profile
Braze or comparable CRM	PROVEN	Braze experience in telecom lifecycle	Prosper / verified profile
CRM audience growth	PROVEN	225% growth to 250K+ subscribers	Bob's Watches / metrics ledger
Cross-functional leadership	PROVEN	Teams, vendors, data, automation	verified profile
End-to-end lifecycle ownership	PLAUSIBLE	Multiple-stage ownership; exact scope needs detail	story bank
Behavioral segmentation	PLAUSIBLE	Transferable CRM work; depth not yet explicit	story bank
Canvas / Liquid mechanics	MISSING	No verified record yet	not found
Push + in-app ownership	MISSING	No verified record yet	not found
Cohort / LTV modeling	RISKY	No externally safe retention/LTV proof	excluded

Decision

Recommendation

Investigate 3 proof gaps

Score breakdown

Lane	93
Mandate	90
Evidence	65
Compensation	100
Freshness	100

Why it ranks first

- ✓ Lifecycle systems + Braze strength align with mandate
- ✓ Remote compensation target met or exceeded
- ✓ Fresh canonical source with recent posting date

Main risk

Advanced Braze mechanics and mobile-channel depth may be below the role's bar.

Pursuit cost

Medium-high · 3 narrative answers + 1–4 hour assessment

Start 8-minute proof check >

Hold

Reject

Correct evidence

No application or Teal changes yet. ⓘ

Evidence in this decision ⓘ

Canonical profile

Updated Jul 16

Metrics ledger

3 sources

Story bank

12 stories

Going careers

going.com/careers

Ashby posting

jobs.ashbyhq.com